

2022/06/10 ~ 2022/06/17

PM stall 관련 실험

Flush instruction 없는 PM DAX write 실험

▼ Cache line 크기 고려한 buffer 문자열 설정

- Cache line은 64 bytes 크기이므로 이를 고려하여 64 bytes보다 큰 주기로 문자열이 반복되게 하거나 아니면 아예 랜덤으로 긴 문자열을 생성시켜 주어야 한다.
- 워크로드 시간 측정 이전에 한 번만 진행해주면 되며, 큰 오버헤드가 없는 것으로 생각되므로 stdlib 헤더파일에 정의된 rand, srand 함수와 time 헤더파일에 정의된 time 함수를 이용해 다음과 같이 난수 생성을 통한 랜덤한 긴 문자열을 생성시켜 준다.

```
int rand_num = 0;
srand((unsigned int)time(NULL));

for(i = 0; i < buf_size_in_byte; i++){
    rand_num = rand()%64;
    buf[i] = '0' + rand_num;
}
```

▼ 파일 생성 코드

```
//Generate 1GB size file for mmap_write_seq workload(The block size can be change, and it will match with the block size which is
#include <sys/types.h>
#include <sys/mman.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <unistd.h>
#include <stdlib.h>
#include <stdio.h>
#include <sys/time.h>
#include <string.h>
#include <errno.h>

int main(){
    int fd;
    int i;
    char *map;
    int GB_to_KB = 1048576; //1 GB = 1048576 KB
    int KB_to_Byte = 1024; //1 KB = 1024 Bytes

    int buf_size_in_byte = 4096; //4KB
    char buf[4096] = {0, }; //4KB

    int iter_num = (GB_to_KB*KB_to_Byte)/buf_size_in_byte;

    for(i = 0; i < buf_size_in_byte; i++){
        buf[i] = 'w';
    }
    if((fd = open("/pmem/file1", O_RDWR|O_CREAT, 0644)) < 0){
        printf("failed to open file\n");
        return 0;
    }
    printf("Starting 1G file generating\n");
    struct timeval startTime, endTime;
    double diffTime;
    gettimeofday(&startTime, NULL);
    for(i = 0; i < iter_num; i++){
        write(fd, buf, 4096);
    }
    gettimeofday(&endTime, NULL);
    if(endTime.tv_usec < startTime.tv_usec){
        endTime.tv_usec += 1000000;
        endTime.tv_sec -= 1;
    }
    printf("The end of file generating operation : %ld.%ld\n", endTime.tv_sec - startTime.tv_sec, endTime.tv_usec - startTime.tv_us);
    close(fd);
    return 0;
}
```

▼ write 워크로드 코드

총 20GB를 write하며, memcpy를 통해 buffer 크기만큼 sequential하게 1GB를 write하는 과정을 20번 반복한다

```
//Workload for writing to PM as DAX FS mode in sequential blocks(the amount of sequential write size is restricted due to the limit of the kernel)
#include <sys/types.h>
#include <sys/mman.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <unistd.h>
#include <stdlib.h>
#include <stdio.h>
#include <sys/time.h>
#include <time.h>
#include <string.h>
#include <stdint.h>
#include <immintrin.h>

int main(){
    int fd;
    int i,j,k;
    char *map;
    int n_GB = 20;
    int GB_to_KB = 1048576; //1 GB = 1048576 KB
    int KB_to_Byte = 1024; //1 KB = 1024 Bytes
    int flush_iter = 4096/64;

    int buf_size_in_byte = 4096; //4KB:4096Bytes -> write in bytes
    char buf[4096] = {0, }; //4KB

    int iter_num = (GB_to_KB*KB_to_Byte)/buf_size_in_byte;
    int rand_num = 0;
    srand((unsigned int)time(NULL));

    for(i = 0; i < buf_size_in_byte; i++){
        rand_num = rand()%64;
        buf[i] = '0' + rand_num;
    }
    if((fd = open("/pmem/file1", O_RDWR|O_CREAT,0644)) < 0){
        printf("failed to open file\n");
        return 0;
    }

    //map = mmap(NULL, buf_size_in_byte, PROT_READ | PROT_WRITE, MAP_SHARED_VALIDATE|MAP_SYNC, fd, 0);
    map = mmap(NULL, GB_to_KB*KB_to_Byte, PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0); //1GB mapping

    //printf("Starting mmap %dGB writing operation1\n", n_GB);
    struct timeval startTime, endTime, syncTime_before, syncTime_after, unmapTime_before, unmapTime_after;
    double diffTime, syncTime, unmapTime;
    gettimeofday(&startTime, NULL);

    for(i=0; i<n_GB; i++){
        for(j=0; j<iter_num; j++){ //1GB sequential write
            char *map_j = map + (buf_size_in_byte*j);
            memcpy(map_j, buf, buf_size_in_byte);
        }
    }

    gettimeofday(&endTime, NULL);
    if(endTime.tv_usec < startTime.tv_usec){
        endTime.tv_usec += 1000000;
        endTime.tv_sec -= 1;
    }
    //printf("The end of mmap write operation1 : %ld.%ld\n", endTime.tv_sec - startTime.tv_sec, endTime.tv_usec - startTime.tv_usec);
    printf("%ld.%ld\n", endTime.tv_sec - startTime.tv_sec, endTime.tv_usec - startTime.tv_usec);
    gettimeofday(&syncTime_before, NULL);
    msync(map, buf_size_in_byte, MS_SYNC);
    munmap(map, buf_size_in_byte);
    close(fd);
    return 0;
}
```

▼ 실험 환경

- Cache size & NUMA information

L1d cache:	1 MiB
L1i cache:	1 MiB
L2 cache:	32 MiB
L3 cache:	44 MiB
NUMA node0 CPU(s):	0,2,4,6,8,10,12,14,16,18,20,22,24,26,28,30,32,34,36,38,40,42,44,46,48,50,52,54,56,58,60,62
NUMA node1 CPU(s):	1,3,5,7,9,11,13,15,17,19,21,23,25,27,29,31,33,35,37,39,41,43,45,47,49,51,53,55,57,59,61,63

- PM0로 DAX write 진행
- jbd_2는 노드0의 8번 코어에 배치 (저널링 과정에서의 remote access 방지)
- 코어의 governor는 전부 performance mode로 설정

▼ 실험 과정

- PM0를 ext4로 포맷 후 DAX로 mount
- PM0에 1GB 크기의 파일 생성 (혹시 모를 요인을 배제하기 위해 파일 생성 프로세스는 노드0의 2번 코어에 배치)

```
taskset -c 2 ./ (파일 생성 실행파일)
```

- jbd_2의 PID를 알아내어 노드0의 8번 코어에 배치

```
ps -aux | grep jbd2 | grep pmem | grep -o '[0-9]*' | head -1
```

```
taskset -cp 8 (jbd_2 PID)
```

- cache drop

```
echo 3 > /proc/sys/vm/drop_caches
```

- 20GB write 워크로드 실행 및 perf 프로파일링

```
perf stat -B -e cache-references,cache-misses,cycles,instructions,branches,faults,migrations,L1-dcache-load-misses,L1-dcache-loads,L1-dcache-stores,L1-icache-load-misses,LLC-loads,LLC-load-misses,LLC-stores,LLC-store-misses,LLC-prefetches,cycle_activity.stalls_l1d_miss,cycle_activity.stalls_l2_miss,cycle_activity.stalls_l3_miss,cycle_activity.stalls_total
taskset -c 2 ./ (워크로드 실행파일)
```

- cf) 편의를 위해 단계 1~4까지를 하나의 bash script로 구성

```
#!/bin/bash

rm /pmem/file1;
umount /pmem;
yes | mkfs.ext4 /dev/pmem0;
mount -o dax /dev/pmem0 /pmem;
taskset -c 2 ./mmap_generate_seq1;
JBDPID=$(ps -aux | grep jbd2 | grep pmem | grep -o '[0-9]*' | head -1)
BINDJBD='taskset -cp 8'
$BINDJBD $JBDPID
echo 3 > /proc/sys/vm/drop_caches;
```

- cf) L1d cache 크기가 1MiB이므로 buffer 크기(한 번에 memcpy로 내리는 크기)는 1MB 보다 작은 크기/1MB/1MB보다 큰 크기, 이렇게 3가지 종류로 설정함
 - 서버 사양인 L1d cache의 크기인 1MB보다 작은 4KB의 random 문자열을 반복적으로 sequential하게 write
 - 서버 사양인 L1d cache의 크기인 1MB보다 작은 256KB의 random 문자열을 반복적으로 sequential하게 write
 - 서버 사양인 L1d cache의 크기인 1MB와 동일한 1MB의 random 문자열을 반복적으로 sequential하게 write
 - 서버 사양인 L1d cache의 크기인 1MB보다 큰 2MB의 random 문자열을 반복적으로 sequential하게 write

▼ 실험 결과

- 4KB 단위 1GB sequential write

▼ perf 결과

```

11.434726

Performance counter stats for 'taskset -c 2 ./mmap_write_seq_4kb':

   339814889      cache-references                      (23.51%)
   333880326      cache-misses                      # 98.254 % of all cache refs (23.54%)
  31896953478      cycles                          (23.58%)
  2674504422      instructions                      # 0.08 insn per cycle (29.49%)
   303288010      branches                          (29.49%)
       31282      faults
           1      migrations
   342595080      L1-dcache-load-misses          # 41.62% of all L1-dcache accesses (29.48%)
   823118476      L1-dcache-loads                      (29.45%)
   731544297      L1-dcache-stores                      (29.42%)
   5462402       L1-icache-load-misses          (23.51%)
   834583        LLC-loads                      (23.51%)
   148972        LLC-load-misses              # 17.85% of all LL-cache accesses (23.51%)
  232298633      LLC-stores                      (11.75%)
       29270      LLC-store-misses          (11.75%)
<not supported> LLC-prefetches
   55178094      cycle_activity.stalls_l1d_miss          (17.63%)
   33662157      cycle_activity.stalls_l2_miss          (23.50%)
   10797535      cycle_activity.stalls_l3_miss          (23.50%)
  30672980126      cycle_activity.stalls_total          (23.50%)

11.440363582 seconds time elapsed

11.336155000 seconds user
 0.099966000 seconds sys

```

- 256KB 단위 1GB sequential write

▼ perf 결과

```

5.354500

Performance counter stats for 'taskset -c 2 ./mmap_write_seq_256kb':

   17936522      cache-references                      (23.32%)
   17519768      cache-misses                      # 97.677 % of all cache refs (23.44%)
 14932138728      cycles                          (23.51%)
 330107533       instructions                      # 0.02 insn per cycle (29.41%)
   94096168      branches                          (29.48%)
       31346      faults
           1      migrations
   677315541      L1-dcache-load-misses          # 768.35% of all L1-dcache accesses (29.54%)
   88152043      L1-dcache-loads                      (29.49%)
   39005446      L1-dcache-stores                      (29.49%)
   4308047       L1-icache-load-misses          (23.59%)
   352816        LLC-loads                      (23.59%)
   250874        LLC-load-misses              # 71.11% of all LL-cache accesses (23.60%)
  16908624      LLC-stores                      (11.80%)
       11486      LLC-store-misses          (11.80%)
<not supported> LLC-prefetches
  4038482515      cycle_activity.stalls_l1d_miss          (17.65%)
   35999872      cycle_activity.stalls_l2_miss          (23.48%)
   24963225      cycle_activity.stalls_l3_miss          (23.40%)
  13257344433      cycle_activity.stalls_total          (23.33%)

5.362095405 seconds time elapsed

5.289539000 seconds user
 0.067968000 seconds sys

```

```

7.519034

Performance counter stats for 'taskset -c 2 ./mmap_write_seq_256kb':

    70804292      cache-references                      (23.40%)
    68349742      cache-misses                        # 96.533 % of all cache refs (23.45%)
   20969238127    cycles                             (23.50%)
    316511491     instructions                       # 0.02 insn per cycle (29.40%)
    89758404      branches                           (29.46%)
    31343          faults
     1            migrations
    678553234     L1-dcache-load-misses              # 786.40% of all L1-dcache accesses (29.50%)
    86285590      L1-dcache-loads                    (29.51%)
    38364045      L1-dcache-stores                    (29.51%)
    4173573       L1-icache-load-misses              (23.61%)
    575885        LLC-loads                         (23.61%)
    253335        LLC-load-misses                   # 43.99% of all LL-cache accesses (23.61%)
    67157752      LLC-stores                         (11.79%)
    23217         LLC-store-misses                 (11.74%)
<not supported> LLC-prefetches
   2386388634     cycle_activity.stalls_l1d_miss      (17.59%)
    43413708      cycle_activity.stalls_l2_miss      (23.44%)
    24353926      cycle_activity.stalls_l3_miss      (23.39%)
   19194035252    cycle_activity.stalls_total        (23.39%)

    7.527847589 seconds time elapsed

    7.415566000 seconds user
    0.107993000 seconds sys

```

```

10.536647

Performance counter stats for 'taskset -c 2 ./mmap_write_seq_256kb':

   192246383      cache-references                      (23.54%)
   187986992      cache-misses                        # 97.784 % of all cache refs (23.54%)
  29393297902     cycles                             (23.54%)
   352189810      instructions                       # 0.01 insn per cycle (29.42%)
   98532867       branches                           (29.42%)
    31348          faults
     1            migrations
   678261306     L1-dcache-load-misses              # 790.64% of all L1-dcache accesses (29.41%)
   85786246       L1-dcache-loads                    (29.41%)
   38467157       L1-dcache-stores                    (29.41%)
   4359212        L1-icache-load-misses              (23.53%)
    754799        LLC-loads                         (23.53%)
    208057        LLC-load-misses                   # 27.56% of all LL-cache accesses (23.53%)
   183295889      LLC-stores                         (11.76%)
    26124         LLC-store-misses                 (11.76%)
<not supported> LLC-prefetches
   1872418182     cycle_activity.stalls_l1d_miss      (17.64%)
    36016608      cycle_activity.stalls_l2_miss      (23.53%)
    16949434      cycle_activity.stalls_l3_miss      (23.53%)
   27683372443    cycle_activity.stalls_total        (23.53%)

   10.545675977 seconds time elapsed

   10.445137000 seconds user
    0.096010000 seconds sys

```

- 1MB 단위 1GB sequential write
 - ▼ perf 결과

```

6.525489

Performance counter stats for 'taskset -c 2 ./mmap_write_seq_1mb':

    195930096      cache-references              (23.55%)
    32870444      cache-misses                 # 16.777 % of all cache refs (23.61%)
    18223909764    cycles                       (23.66%)
    529476740      instructions                 # 0.03  insn per cycle       (29.53%)
    126819035      branches                     (29.53%)
    31535          faults
    1            migrations
    676044650      L1-dcache-load-misses        # 604.10% of all L1-dcache accesses (29.47%)
    111909316      L1-dcache-loads              (29.41%)
    38399752       L1-dcache-stores             (29.36%)
    4153566        L1-icache-load-misses        (23.49%)
    61538123       LLC-loads                    (23.49%)
    238932         LLC-load-misses             # 0.39% of all LL-cache accesses (23.49%)
    31889961       LLC-stores                   (11.74%)
    13383          LLC-store-misses            (11.75%)
<not supported>  LLC-prefetches
    5936264477     cycle_activity.stalls_l1d_miss (17.62%)
    2111179890     cycle_activity.stalls_l2_miss (23.49%)
    19045782       cycle_activity.stalls_l3_miss (23.49%)
    16447293005    cycle_activity.stalls_total   (23.49%)

    6.543570829 seconds time elapsed

    6.467334000 seconds user
    0.072037000 seconds sys

```

```

10.639645

Performance counter stats for 'taskset -c 2 ./mmap_write_seq_1mb':

    492790744      cache-references              (23.47%)
    180090282      cache-misses                 # 36.545 % of all cache refs (23.51%)
    29709815582    cycles                       (23.54%)
    584115331      instructions                 # 0.02  insn per cycle       (29.44%)
    142141615      branches                     (29.48%)
    31536          faults
    1            migrations
    676776687      L1-dcache-load-misses        # 605.07% of all L1-dcache accesses (29.47%)
    111851277      L1-dcache-loads              (29.47%)
    36584558       L1-dcache-stores             (29.47%)
    3895522        L1-icache-load-misses        (23.58%)
    62339981       LLC-loads                    (23.58%)
    238142         LLC-load-misses             # 0.38% of all LL-cache accesses (23.58%)
    176949462     LLC-stores                   (11.76%)
    17968          LLC-store-misses            (11.73%)
<not supported>  LLC-prefetches
    3505554040     cycle_activity.stalls_l1d_miss (17.58%)
    1808797756     cycle_activity.stalls_l2_miss (23.44%)
    16269216       cycle_activity.stalls_l3_miss (23.43%)
    27940695633    cycle_activity.stalls_total   (23.43%)

    10.657553717 seconds time elapsed

    10.533405000 seconds user
    0.119970000 seconds sys

```

- 2MB 단위 1GB sequential write
 - ▼ perf 결과

```

7.133169

Performance counter stats for 'taskset -c 2 ./mmap_write_seq_2mb':

   367911270      cache-references                      (23.48%)
   29634778       cache-misses                        #    8.055 % of all cache refs (23.53%)
  19955469544     cycles                                (23.59%)
   599393221     instructions                        #    0.03  insn per cycle     (29.51%)
  160128369      branches                                (29.57%)
    31792        faults
         1        migrations
   670126020     L1-dcache-load-misses      #  317.81% of all L1-dcache accesses (29.59%)
  210854678      L1-dcache-loads                      (29.53%)
   87321598      L1-dcache-stores                      (29.48%)
   3435227       L1-icache-load-misses                (23.50%)
  156059956      LLC-loads                          (23.47%)
   296645        LLC-load-misses                    #    0.19% of all LL-cache accesses (23.47%)
  28492617       LLC-stores                          (11.73%)
    9254         LLC-store-misses                    (11.73%)
<not supported>  LLC-prefetches
   9219383550    cycle_activity.stalls_l1d_miss          (17.60%)
   4887131977    cycle_activity.stalls_l2_miss          (23.47%)
   20780418      cycle_activity.stalls_l3_miss          (23.47%)
  18124563518    cycle_activity.stalls_total          (23.47%)

   7.162495314 seconds time elapsed

   7.071309000 seconds user
   0.087991000 seconds sys

```

```

10.520876

Performance counter stats for 'taskset -c 2 ./mmap_write_seq_2mb':

   512192926      cache-references                      (23.55%)
  170419521       cache-misses                        #   33.273 % of all cache refs (23.58%)
 29410590599     cycles                                (23.58%)
   681044634     instructions                        #    0.02  insn per cycle     (29.46%)
  184112231      branches                                (29.46%)
    31794        faults
         1        migrations
   675021802     L1-dcache-load-misses      #  331.18% of all L1-dcache accesses (29.42%)
  203822656      L1-dcache-loads                      (29.39%)
   82359338      L1-dcache-stores                      (29.39%)
   4092465       L1-icache-load-misses                (23.51%)
   70091403      LLC-loads                          (23.51%)
   240579        LLC-load-misses                    #    0.34% of all LL-cache accesses (23.51%)
  167422497      LLC-stores                          (11.76%)
    7612         LLC-store-misses                    (11.76%)
<not supported>  LLC-prefetches
  3889698229     cycle_activity.stalls_l1d_miss          (17.64%)
  2070170699     cycle_activity.stalls_l2_miss          (23.51%)
   17064804      cycle_activity.stalls_l3_miss          (23.51%)
  27646875735    cycle_activity.stalls_total          (23.51%)

  10.551004758 seconds time elapsed

  10.474609000 seconds user
   0.072017000 seconds sys

```

▼ 문제점

- 동일한 조건으로 실험을 진행했을 때, 결과값이 다른 경우가 존재함 (jbd_2를 동일 노드 특정 코어에 고정시켰을 때도 결과값이 다른 경우가 발생함을 확인)
- 단, 4KB 랜덤 문자열 버퍼로 구성되어 4KB씩 write를 진행하는 워크로드의 경우 4회 실험 전부 동일 결과를 보여줌
- 아래의 두 perf 프로파일링 결과는 동일 조건에서 실행한 2MB 랜덤 문자열 버퍼로 구성된 1GB sequential PM DAX write without flush instruction 워크로드 실험 결과임

```

7.83950

Performance counter stats for 'taskset -c 2 ./mmap_write_seq_2mb':

   367045662      cache-references              (23.01%)
   29087229       cache-misses                #    7.925 % of all cache refs (30.72%)
  19821736833     cycles                  (30.78%)
   544606338     instructions                #    0.03  insn per cycle     (38.48%)
  145962544      branches                  (38.54%)
    31792        faults
         1        migrations
   676378366     L1-dcache-load-misses      #  399.30% of all L1-dcache accesses (38.54%)
  169391093      L1-dcache-loads              (38.54%)
   69759779      L1-dcache-stores              (38.54%)
   4655427       L1-icache-load-misses        (30.83%)
  156919113      LLC-loads                  (30.80%)
   251706        LLC-load-misses            #    0.16% of all LL-cache accesses (30.74%)
  28109500       LLC-stores                  (15.30%)
    8821         LLC-store-misses            (15.30%)
<not supported>  LLC-prefetches

   7.114172163 seconds time elapsed

   7.005999000 seconds user
   0.103970000 seconds sys

```

```

9.830712

Performance counter stats for 'taskset -c 2 ./mmap_write_seq_2mb':

   485695689      cache-references              (23.02%)
  144601495       cache-misses                #   29.772 % of all cache refs (30.73%)
 27491330676     cycles                  (30.77%)
   506069044     instructions                #    0.02  insn per cycle     (38.48%)
  133547768      branches                  (38.52%)
    31795        faults
         1        migrations
   676775935     L1-dcache-load-misses      #  376.66% of all L1-dcache accesses (38.55%)
  179680039      L1-dcache-loads              (38.55%)
   76999925      L1-dcache-stores              (38.55%)
   4034046       L1-icache-load-misses        (30.80%)
   59078614      LLC-loads                  (30.76%)
   233596        LLC-load-misses            #    0.40% of all LL-cache accesses (30.72%)
  142429585      LLC-stores                  (15.34%)
    16467        LLC-store-misses            (15.34%)
<not supported>  LLC-prefetches

   9.860411210 seconds time elapsed

   9.736797000 seconds user
   0.120009000 seconds sys

```

PM DAX read 실험(Read이므로 flush instruction은 기본적으로 없음)

▼ Buffer 문자열 설정

- 실험 구성이 4KB, 256KB, 1MB, 2MB 단위로 1GB를 sequential하게 읽어오기 때문에 적어도 2MB보다 큰 랜덤 문자열로 buffer를 구성해주는 것이 좋다.
- 따라서 위의 PM DAX write 실험의 random 문자열 생성 코드를 이용하여 4MB 크기의 buffer에 랜덤 문자열을 생성해주었다.

▼ 파일 생성 코드

- PM DAX write 실험과 다르게 cache line으로의 흡수를 막기 위해 4MB 단위의 랜덤 문자열로 write를 진행하여 1GB 크기의 파일을 생성하도록 하였다.

```

//Generate 1GB size file for mmap_write_seq workload(The block size can be change, and it will match with the block size which
#include <sys/types.h>
#include <sys/mman.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <unistd.h>
#include <stdlib.h>
#include <stdio.h>
#include <sys/time.h>
#include <time.h>
#include <string.h>
#include <errno.h>

```



```

int main(){
    int fd;
    int i;
    char *map;
    int GB_to_KB = 1048576; //1 GB = 1048576 KB
    int KB_to_Byte = 1024; //1 KB = 1024 Bytes

    int buf_size_in_byte = 4096*1024; //2MB
    char buf[4096*1024] = {0, }; //2MB

    int iter_num = (GB_to_KB*KB_to_Byte)/buf_size_in_byte;

    srand((unsigned int)time(NULL));
    int rand_num = 0;
    for(i = 0; i < buf_size_in_byte; i++){
        rand_num = rand()%64;
        buf[i] = '0' + rand_num;
    }
    if((fd = open("/pmem/file1", O_RDWR|O_CREAT,0644)) < 0){
        printf("failed to open file\n");
        return 0;
    }
    printf("Starting 1G file generating1\n");
    struct timeval startTime, endTime;
    double diffTime;
    gettimeofday(&startTime, NULL);
    for(i = 0; i < iter_num; i++){
        write(fd, buf, buf_size_in_byte);
    }
    gettimeofday(&endTime, NULL);
    if(endTime.tv_usec < startTime.tv_usec){
        endTime.tv_usec += 1000000;
        endTime.tv_sec -= 1;
    }
    printf("The end of file generating operation1 : %ld.%ld\n", endTime.tv_sec - startTime.tv_sec, endTime.tv_usec - startTime.tv_usec);
    close(fd);
    return 0;
}

```

▼ read 워크로드 코드

memcpy read 속도가 write에 비해 약 2배 정도 빠른 것을 고려하여 sequential하게 1GB 파일을 총 40번 읽어들이 전체 40GB를 read하는 워크로드를 구성하였다.

```

//Workload for writing to PM as DAX FS mode in sequential blocks(the amount of sequential write size is restricted due to the limit)
#include <sys/types.h>
#include <sys/mman.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <unistd.h>
#include <stdlib.h>
#include <stdio.h>
#include <sys/time.h>
#include <time.h>
#include <string.h>
#include <stdint.h>
#include <immintrin.h>

int main(){
    int fd;
    int i,j,k;
    char *map;
    int n_GB = 40;
    int GB_to_KB = 1048576; //1 GB = 1048576 KB
    int KB_to_Byte = 1024; //1 KB = 1024 Bytes
    int flush_iter = 4096/64;

    int buf_size_in_byte = 4096; //4KB:4096Bytes -> write in bytes
    char buf[4096] = {0, }; //4KB

    int iter_num = (GB_to_KB*KB_to_Byte)/buf_size_in_byte;
    if((fd = open("/pmem/file1", O_RDWR|O_CREAT,0644)) < 0){
        printf("failed to open file\n");
        return 0;
    }

    //map = mmap(NULL, buf_size_in_byte, PROT_READ | PROT_WRITE, MAP_SHARED_VALIDATE|MAP_SYNC, fd, 0);
    map = mmap(NULL, GB_to_KB*KB_to_Byte, PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0); //1GB mapping

    //printf("Starting mmap %dGB writing operation1\n", n_GB);
    struct timeval startTime, endTime, syncTime_before, syncTime_after, unmapTime_before, unmapTime_after;
    double diffTime, syncTime, unmapTime;

```

```

gettimeofday(&startTime, NULL);

for(i=0; i<n_GB; i++){
    for(j=0; j<iter_num; j++){ //1GB sequential write
        char *map_j = map + (buf_size_in_byte*j);
        memcpy(buf, map_j, buf_size_in_byte);
    }
}

gettimeofday(&endTime, NULL);
if(endTime.tv_usec < startTime.tv_usec){
    endTime.tv_usec += 1000000;
    endTime.tv_sec -= 1;
}
//printf("The end of mmap write operation1 : %ld.%ld\n", endTime.tv_sec - startTime.tv_sec, endTime.tv_usec - startTime.tv_usec);
printf("%ld.%ld\n", endTime.tv_sec - startTime.tv_sec, endTime.tv_usec - startTime.tv_usec);
gettimeofday(&syncTime_before, NULL);
msync(map, buf_size_in_byte, MS_SYNC);
munmap(map, buf_size_in_byte);
close(fd);
return 0;
}

```

▼ 실험 환경

- Cache size & NUMA information

```

L1d cache:      1 MiB
L1i cache:      1 MiB
L2 cache:       32 MiB
L3 cache:       44 MiB
NUMA node0 CPU(s): 0,2,4,6,8,10,12,14,16,18,20,22,24,26,28,30,32,34,36,38,40,42,44,46,48,50,52,54,56,58,60,62
NUMA node1 CPU(s): 1,3,5,7,9,11,13,15,17,19,21,23,25,27,29,31,33,35,37,39,41,43,45,47,49,51,53,55,57,59,61,63

```

- PM0로 DAX read 진행
- jbd_2는 노드0의 8번 코어에 배치 (저널링 과정에서의 remote access 방지)
→ read 실험이어서 불필요한 과정이지만, 혹시 모를 요인 배제를 위해 수행
- 코어의 governor는 전부 performance mode로 설정

▼ 실험 과정

PM DAX write 실험과 동일

▼ 실험 결과

- 4KB 단위 1GB sequential read
 - ▼ perf 결과

```

11.976557

Performance counter stats for 'taskset -c 2 ./mmap_read_seq_4kb':

   674675978      cache-references              (23.54%)
   664671047      cache-misses                # 98.517 % of all cache refs (23.57%)
  33406992644      cycles                      (23.59%)
  4816170188      instructions                # 0.14 insn per cycle       (29.47%)
   437527844      branches                    (29.47%)
         626      faults
          1      migrations
   672641661      L1-dcache-load-misses      # 45.59% of all L1-dcache accesses (29.44%)
  1475496532      L1-dcache-loads            (29.40%)
  1388580799      L1-dcache-stores            (29.39%)
   1083688      L1-icache-load-misses      (23.51%)
   556651299      LLC-loads                  (23.51%)
   197046      LLC-load-misses          # 0.04% of all LL-cache accesses (23.51%)
   223072      LLC-stores                (11.76%)
   34341      LLC-store-misses        (11.76%)
<not supported>  LLC-prefetches
  30381979393      cycle_activity.stalls_l1d_miss (17.63%)
  30379477030      cycle_activity.stalls_l2_miss (23.51%)
  30336020958      cycle_activity.stalls_l3_miss (23.51%)
  30405952157      cycle_activity.stalls_total  (23.51%)

11.979840465 seconds time elapsed

11.969299000 seconds user
0.007998000 seconds sys

```

- 256KB 단위 1GB sequential read

▼ perf 결과

```

10.880918

Performance counter stats for 'taskset -c 2 ./mmap_read_seq_256kb':

   673058896      cache-references              (23.53%)
   668648782      cache-misses                # 99.345 % of all cache refs (23.54%)
  30343869622      cycles                      (23.54%)
   41999740      instructions                # 0.00 insn per cycle       (29.42%)
   9340891      branches                    (29.42%)
         686      faults
          1      migrations
  1347687756      L1-dcache-load-misses      # 15474.24% of all L1-dcache accesses (29.42%)
   8709233      L1-dcache-loads            (29.41%)
   4872508      L1-dcache-stores            (29.41%)
   604810      L1-icache-load-misses      (23.53%)
   565803808      LLC-loads                  (23.53%)
   186455      LLC-load-misses          # 0.03% of all LL-cache accesses (23.53%)
   134612      LLC-stores                (11.76%)
   13109      LLC-store-misses        (11.76%)
<not supported>  LLC-prefetches
  26954760566      cycle_activity.stalls_l1d_miss (17.64%)
  26939135825      cycle_activity.stalls_l2_miss (23.53%)
  26906697983      cycle_activity.stalls_l3_miss (23.53%)
  26978778381      cycle_activity.stalls_total  (23.53%)

10.883598043 seconds time elapsed

10.877731000 seconds user
0.004000000 seconds sys

```

- 1MB 단위 1GB sequential read

▼ perf 결과

```

11.352384
Performance counter stats for 'taskset -c 2 ./mmap_read_seq_1mb':

    1013376625      cache-references              (23.51%)
    668849746      cache-misses                 # 66.002 % of all cache refs (23.51%)
    31665535463     cycles                       (23.51%)
    42682576       instructions                 # 0.00 insn per cycle         (29.40%)
    8502111        branches                    (29.41%)
    878            faults
    1              migrations
    1343655847      L1-dcache-load-misses       # 17670.64% of all L1-dcache accesses (29.42%)
    7603889        L1-dcache-loads             (29.42%)
    4596394        L1-dcache-stores            (29.42%)
    657269         L1-icache-load-misses       (23.53%)
    529662691      LLC-loads                   (23.53%)
    183364         LLC-load-misses            # 0.03% of all LL-cache accesses (23.53%)
    144676836      LLC-stores                  (11.77%)
    11328          LLC-store-misses         (11.77%)
<not supported>   LLC-prefetches
    28244472638     cycle_activity.stalls_l1d_miss (17.65%)
    28168276754     cycle_activity.stalls_l2_miss (23.53%)
    28135149042     cycle_activity.stalls_l3_miss (23.53%)
    28292370124     cycle_activity.stalls_total  (23.52%)

    11.356707998 seconds time elapsed

    11.337656000 seconds user
    0.015996000 seconds sys

```

- 2MB 단위 1GB sequential read

▼ perf 결과

```

11.443904
Performance counter stats for 'taskset -c 2 ./mmap_read_seq_2mb':

    695396141      cache-references              (23.51%)
    666471862      cache-misses                 # 95.841 % of all cache refs (23.55%)
    31924250315     cycles                       (23.58%)
    42929103       instructions                 # 0.00 insn per cycle         (29.49%)
    9123646        branches                    (29.52%)
    1136           faults
    1              migrations
    1343503592      L1-dcache-load-misses       # 15825.94% of all L1-dcache accesses (29.52%)
    8489249        L1-dcache-loads             (29.49%)
    5143820        L1-dcache-stores            (29.45%)
    660491         L1-icache-load-misses       (23.51%)
    530326788      LLC-loads                   (23.48%)
    185280         LLC-load-misses            # 0.03% of all LL-cache accesses (23.48%)
    15283792      LLC-stores                  (11.74%)
    12791          LLC-store-misses         (11.74%)
<not supported>   LLC-prefetches
    28496757413     cycle_activity.stalls_l1d_miss (17.61%)
    28286559049     cycle_activity.stalls_l2_miss (23.48%)
    28247008374     cycle_activity.stalls_l3_miss (23.48%)
    28564343645     cycle_activity.stalls_total  (23.48%)

    11.448833524 seconds time elapsed

    11.437904000 seconds user
    0.007998000 seconds sys

```

다른 방향성

turbo mode

- 각 코어별 전력은 제한되어있음 → 만약 PM으로의 I/O가 전력을 덜 소비하면 sibling core에서 돌아가는 task가 더 많은 전력을 소비해서 더 좋은 성능을 낼 수 있는 것 아닐까? → 온도는 HW의 온도 센서가 측정하지만, 전력은 Idle cycle(?)의 비율에 따라 책정되는

데 PM으로의 I/O가 과연 어떤지는 모르겠음

→ 스케줄링으로 연결됨

Noisy neighbor

- 우선 clflush와 다르게 clflushopt와 clwb이 sibling core의 CPU task의 성능을 상당히 저하시키는 현상의 원인은 clflushopt와 clwb이 bus mastering을 하기 때문일수도?(clflushopt와 clwb이 bus mastering을 하는게 확실한 것은 아님. 그냥 가정일 뿐. 확인 필요)
- 지금까지 발견한 사실은 현재 PM DAX I/O stack 상 msync의 마지막 instruction이 clwb인데, 이 clwb이 I/O의 성능은 높일지 몰라도 noisy neighbor가 될 수 있음 (clflush와 다르게 sibling core에 있는 CPU task의 성능을 심각하게 저하시키는 것은 확인된 사실이며, LLC를 공유하는 task들의 성능 또한 저하시킬 여지가 있음 → LLC 공유 관련은 실험 필요)
- 반면, clflush는 sibling core의 CPU task의 성능을 대체로 향상시키지만 I/O task 본인의 성능은 clwb에 비해 상당히 감소됨
- 따라서 우리는 현재 PM I/O stack이 noisy neighbor가 될 수 있음을 발견하였기에 clflush와 clwb 사이의 절충안을 제시하려함

→ Noisy neighbor이므로 sibling core나 LLC를 고려한 스케줄링으로 연결됨

→ clflush와 clwb 사이의 절충안, 즉 flush 및 cache 관련쪽으로 연결됨